

August 24th, 2018

# Refugee Predicting Refugee Migration

---

Prepared by: Jose Pauig



## Abstract

By the end of 2014, the estimate for total internally displaced persons and refugees was sixty million - a staggering amount. If machine learning has a role in aiding the resettlement process it might lie in predicting the future movements of refugees. If we knew in advance where refugees will originate, could the functions governing application, training, and support - the systems which drive successful resettlement - operate more effectively?

IPYNB files accompany this report and include the code which is the basis of this analysis.

## Objective:

1. To build a model which takes as input a country represented by its demographic, economic, health, and trade measures and outputs a decision as to whether the country will generate refugees in a particular year, and;
2. To build a model which takes as input a country represented by its demographic, economic, health and trade measures and outputs a prediction on the number of

refugees generated in a particular year and the destinations (at a state level in the US) of those refugees.

3. To see if different clusters of countries generate different volumes.

## Sourcing and Cleaning our Data

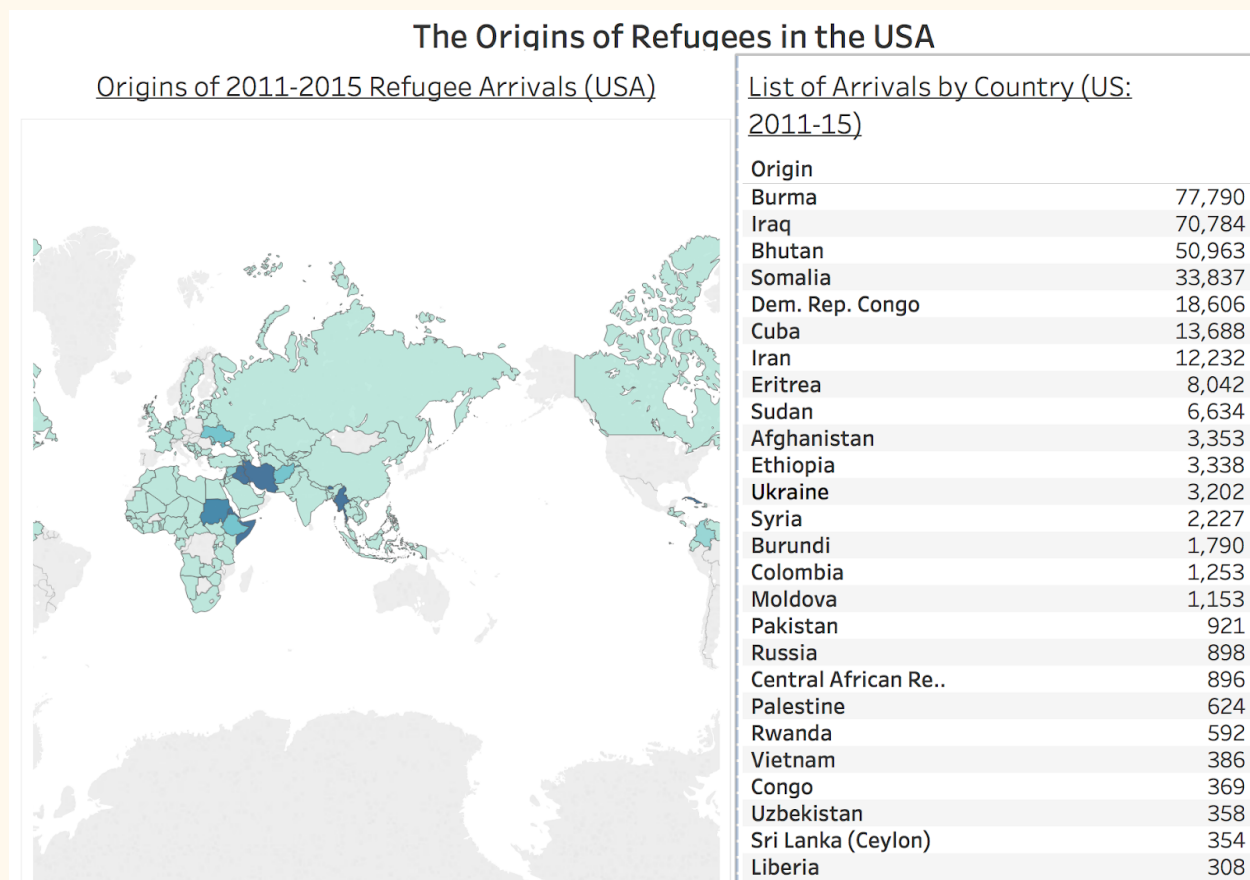
Information representing the demographic, economic, health, work and trade measures of all countries recognized by the UN was sourced from: <http://hdr.undp.org/en/data>. The dependent variables for our models will comprise of this information.

Information about the volume, origins and destinations of refugees arriving in the US for each of the five years ending 2015 were taken from BuzzFeed News' Github repository: <https://github.com/BuzzFeedNews>. It was originally collected from The Department of State's Refugee Processing Center.


As the UN dataset contained missing information - missing countries, fields with missing values - certain countries and fields were discarded. For example any measure for which more than 20 countries were lacking information was removed from the dataset. Where a country or field was missing only a few values, the respective mean replaced the missing data.

## Gaining Familiarity over our Data

Plotting the locations of origin countries that generate refugees, it becomes apparent that the highest volume origins generally cluster in East Africa and in two areas in Asia. Certainly some of the countries which output the greatest volumes of refugees come as no surprise. Burma which struggles with an ethnic conflict involving its Rohingya population, tops Iraq (from which most refugees claim to be escaping ISIS held territory). That Bhutan displaced 51 thousand refugees over our time period may come as a surprise. These countries raise two considerations: (1) There are a small handful of countries which generate the bulk of refugees, and (2) conflicts that are specific to these countries may explain the resulting displacement.



Given that our target variables are strictly positive we might expect a right skewed distribution for arrivals. As pictured below, significantly all countries output a low volume if any at all.




Our predictor variables are as follows:

{'Demography4': 'Population, ages 15\>64 (millions)',  
 'Demography5': 'Population, ages 65 and older (millions)',  
 'Demography6': 'Population, total (millions)',  
 'Demography7': 'Population, under age 5 (millions)',  
 'Demography8': 'Population, urban (%)',  
 'Education2': 'Expected years of schooling (years)',  
 'Education3': 'Education Index',  
 'Education9': 'Mean years of schooling (years)',  
 'Environment3': 'Forest area (% of total land area)',  
 'Equality11': 'Inequality-adjusted life expectancy index',  
 'Equality7': 'Inequality in life expectancy (%)',  
 'Gender1': 'Adolescent birth rate (births per 1,000 women ages 15-19)',  
 'Gender11': 'Labour force participation rate, female (% ages 15 and older)',  
 'Gender12': 'Labour force participation rate, male (% ages 15 and older)',  
 'Gender14': 'Life expectancy at birth, female (years)',  
 'Gender15': 'Life expectancy at birth, male (years)',  
 'Gender16': 'Maternal mortality ratio (deaths per 100,000 live births)',  
 'Gender17': 'Mean years of schooling, female (years)',  
 'Gender18': 'Mean years of schooling, male (years)',  
 'Gender2': 'Estimated gross national income per capita, female (2011 PPP\$)',  
 'Gender21': 'Share of seats in parliament (% held by women)',  
 'Gender22': 'Unemployment rate (total), female to male ratio',  
 'Gender23': 'Youth unemployment rate, female to male ratio',  
 'Gender3': 'Estimated gross national income per capita, male (2011 PPP\$)',  
 'Gender4': 'Expected years of schooling, female (years)',  
 'Gender5': 'Expected years of schooling, male (years)',  
 'Health10': 'Life expectancy Index',  
 'Health13': 'Under-five mortality rate (per 1,000 live births)',  
 'Health6': 'Infant mortality rate (per 1,000 live births)',  
 'Health9': 'Life expectancy at birth (years)',  
 'Income5': 'Gross domestic product (GDP) per capita (2011 PPP \$)',  
 'Income6': 'Gross domestic product (GDP), total (2011 PPP \$ billions)',  
 'Income8': 'Gross national income (GNI) per capita (2011 PPP\$)',  
 'Income9': 'Income Index',  
 'Mobility3': 'Internet users (% of population)',  
 'Mobility4': 'Mobile phone subscriptions (per 100 people)',

'Trade4': 'Foreign direct investment, net inflows (% of GDP)',  
 'Work4': 'Employment to population ratio (% ages 15 and older)',  
 'Work5': 'Labour force participation rate (% ages 15 and older)',  
 'Work8': 'Total unemployment rate (% of labour force)',

Using a Gapminder plot, we can visualize how arrivals change over two dimensions and over time. Included in the accompanying ipynb files is the gapminder file used in this analysis. Below, arrivals are plotted against GDP, life expectancy and year.

### Life Expectancy, GDP and Year vs. Arrivals by Country



## Building a Classifier to Identify Refugee Generating Countries

Using a Random Forest Classifier, we get  $R^2 = 0.9$  on a validation set (which arbitrarily was set to the 2015 data). We can therefore provide a reasonable prediction on whether a country will generate refugees.

## **Building a Regressor to Predict Volumes and Destination States of Refugees from an Origin Country**

Using a Random Forest Regressor - which outperforms a gradient boosting regressor and linear regressor with regularization - we only get  $R^2 = 0.2$ , a relatively weak result. A log transformation on the dependent variable improves  $R^2$  dramatically but abstracts the problem such that it loses its value as measured against our original objective.

Mapping the residuals, we know that we underperform disproportionately on those countries which generate the highest refugees. As we noted (above) the countries generating the greatest volumes are involved in specific conflicts which may not necessarily be represented by measures in our dependent variable. I believe that improving the performance of the model would require additional fields which would do so.

## **Unsupervised Clustering to Identify Groups of Countries**

Using k\_means where  $k=5$ , identifies one small group of 10 countries including China, Armenia, Denmark, Albania, Lebanon, Argentina, El Salvador, Antigua and Barbuda, Burundi, and Albania. These countries do not share a common geography, level of development or any other immediately apparent attribute. Why they do cluster is a similarity in (1) low frequency of Adolescent Birth, and (2) low frequency of maternal mortality.

Though the within group distributions of arrivals for each of the 5 country clusters are not normal (as we expect the distribution is very much right skewed as a result of those countries which do not generate any refugees at all). But excusing the very far left of the distributions, they are somewhat smooth. If we apply an Anova test, we get a significant p-value which indicates that these groups do in fact generate different volumes.